

```
#Use the transformers library from Hugging Face to load a model and generate text from a prompt:
```

```
#"Write a small poem about the ocean". Print the result.
```

```
from transformers import pipeline
```

```

/usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it as secret.
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
  warnings.warn(
config.json: 100% 665/665 [00:00<00:00, 47.3kB/s]
model.safetensors: 100% 548M/548M [00:17<00:00, 48.9MB/s]
generation_config.json: 100% 124/124 [00:00<00:00, 7.19kB/s]
tokenizer_config.json: 100% 26.0/26.0 [00:00<00:00, 1.41kB/s]
vocab.json: 100% 1.04M/1.04M [00:00<00:00, 8.39MB/s]
merges.txt: 100% 456k/456k [00:00<00:00, 19.8MB/s]
tokenizer.json: 100% 1.36M/1.36M [00:00<00:00, 15.3MB/s]
Device set to use cpu
Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncation=True` to explicitly truncate
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
Both `max_new_tokens` (=256) and `max_length` (=50) seem to have been set. `max_new_tokens` will take precedence. Please refer to the doc
Write a small poem about the ocean, "The Ocean" is an adventure book about a child who finds out her parents have died. The kids' teacher
The kids' teacher is a tragic and tragic woman named Mary, and the children have a very different life. The main character (played by a
The kids' teacher is a sad, sad, sad story about the death of her father,
```

```

generator = pipeline("text-generation", model="gpt2")
prompt = "Write a small poem about the ocean"
# Generate text
result = generator(prompt, max_length=50, num_return_sequences=1)
print(result[0]['generated_text'])
```

```

Device set to use cpu
Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncation=True` to explicitly truncate
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
Both `max_new_tokens` (=256) and `max_length` (=50) seem to have been set. `max_new_tokens` will take precedence. Please refer to the doc
Write a small poem about the ocean and its inhabitants.
It's a simple idea, like, the ocean is the ocean of life. You can't just say, "Oh, my God, I want to believe that this is the ocean we're
So there's this whole idea that we're all in this place, and we are all living in different ways. We're all living in different ways. Th
And the ocean is the ocean of life. It's the ocean of life. It's the ocean of life. It's the ocean of life. It's the ocean of life.
And the ocean is the ocean of life. It's the ocean of life. It's the ocean of life.
```